Kernel-Mode NVRAM

Jun 2011

Agenda

- Original Design
- Issues
- New Design
- Code Review
- Test Scenario

Original Design

- User-space library
- Provide APIs for NVRAM operations
- Allocate a user-space memory region to store NVRAM data from Flash, and each user-space program has individual one
- Data is stored in ASCII format
- API:

  nvram_init, nvram_get, nvram_bufget, nvram_set, nvram_bufset,nvram_commit, nvram_close, nvram_clear

Original Design: nvram_init()

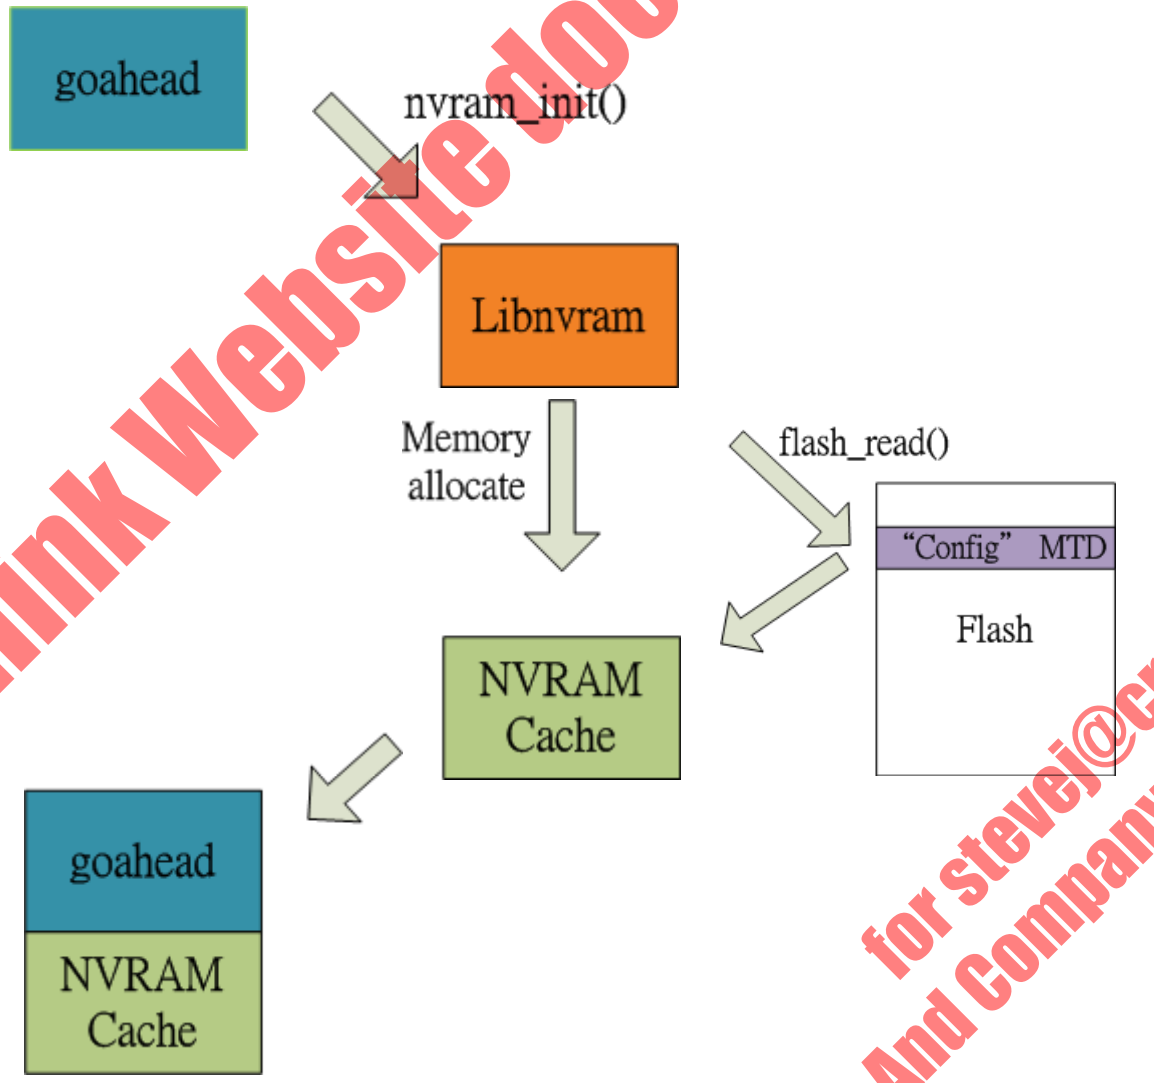

Original Design: nvram_bufget()

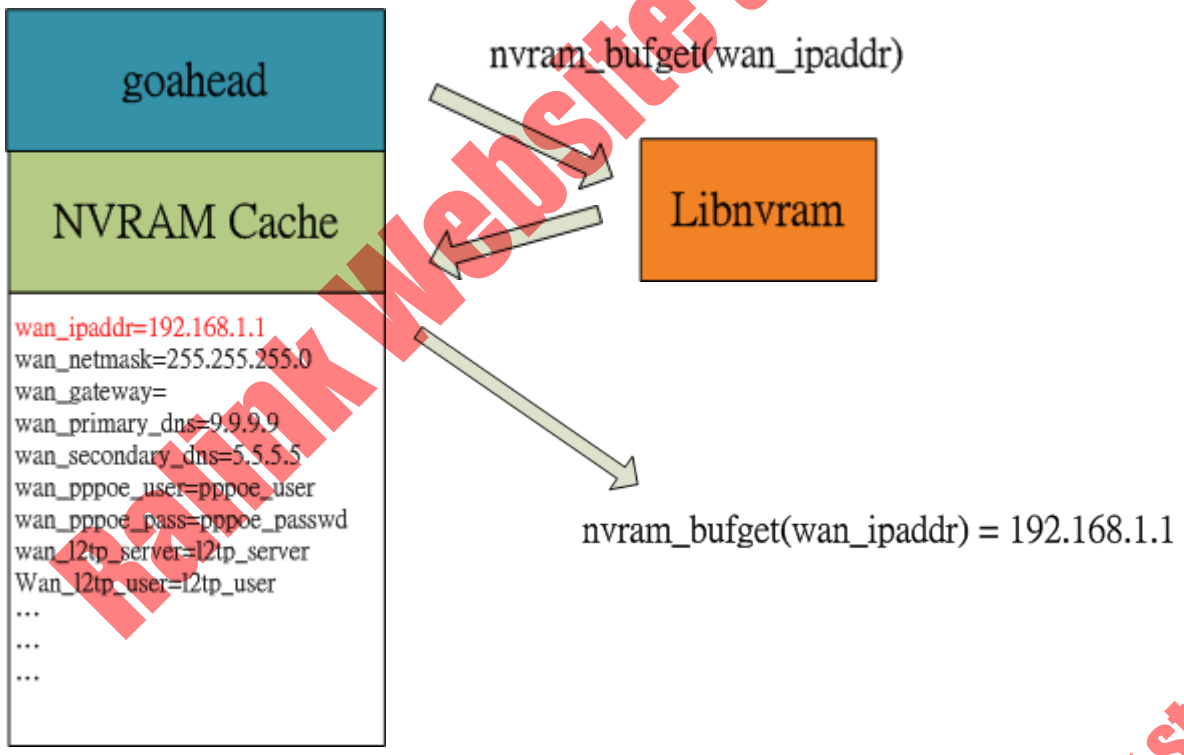

Original Design: nvram_bufset()

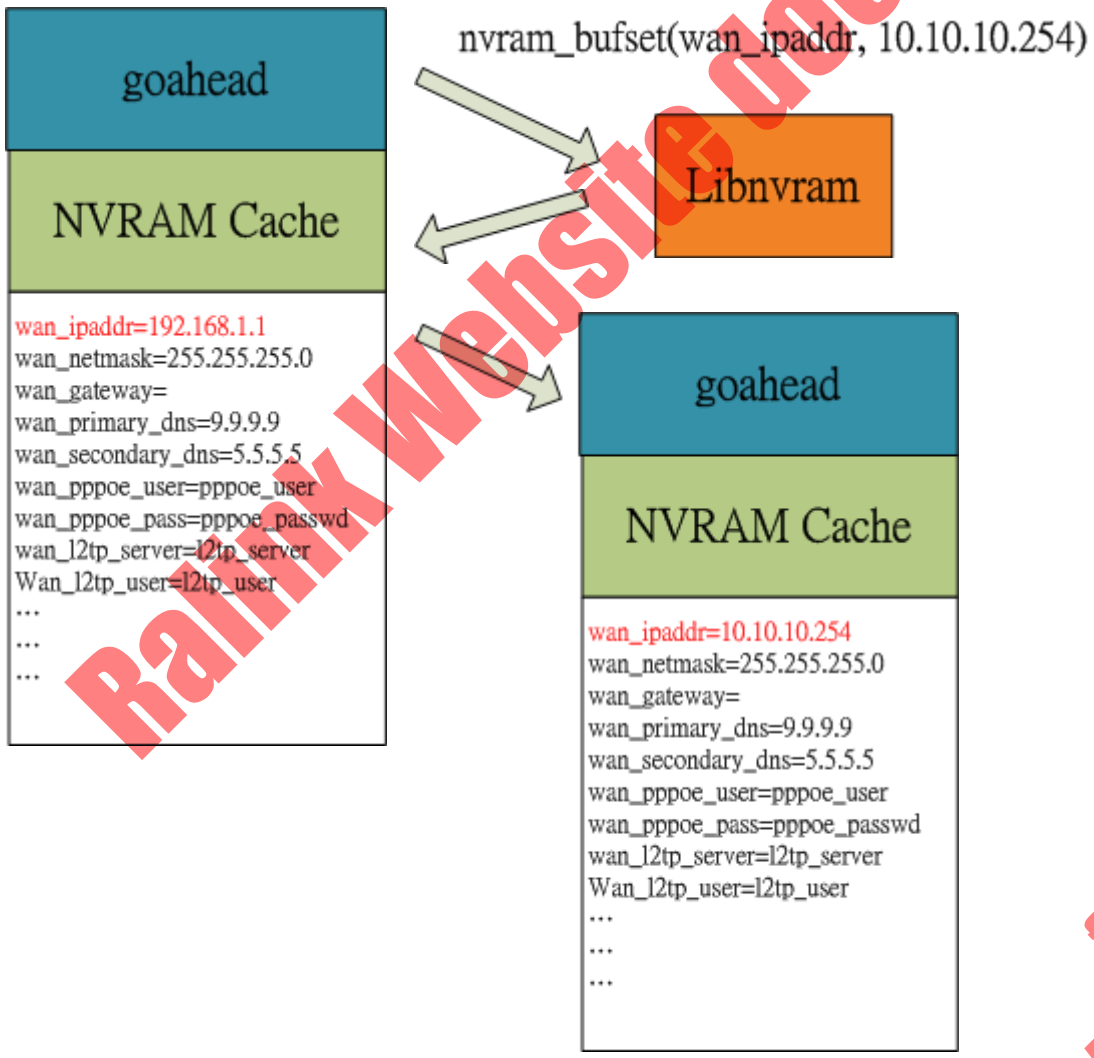

Original Design: nvram_commit()

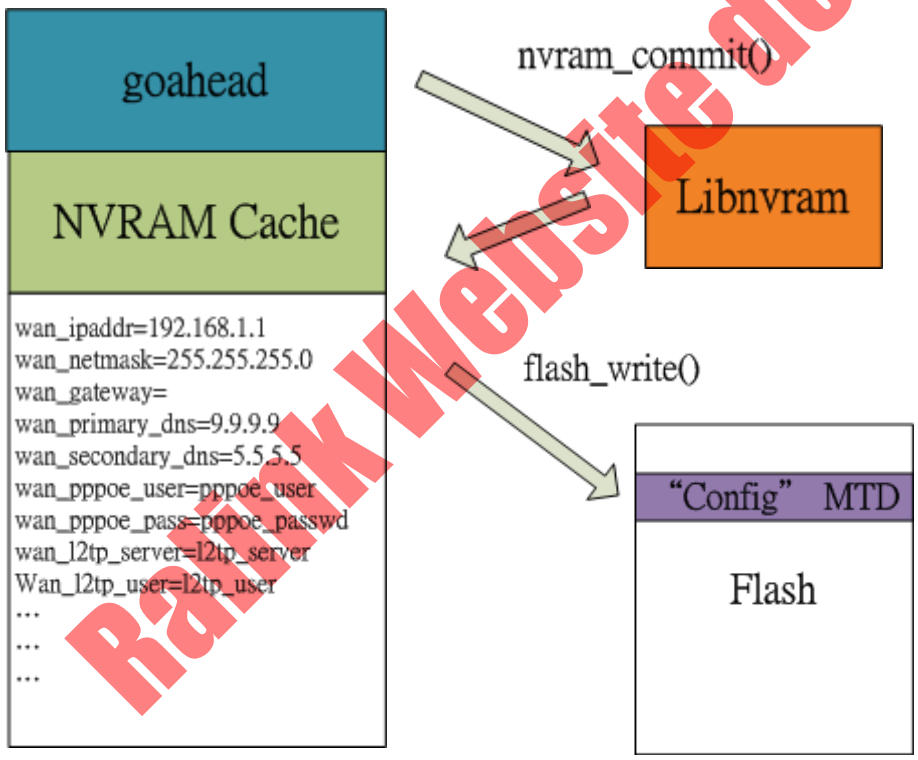

Original Design: nvram_close()

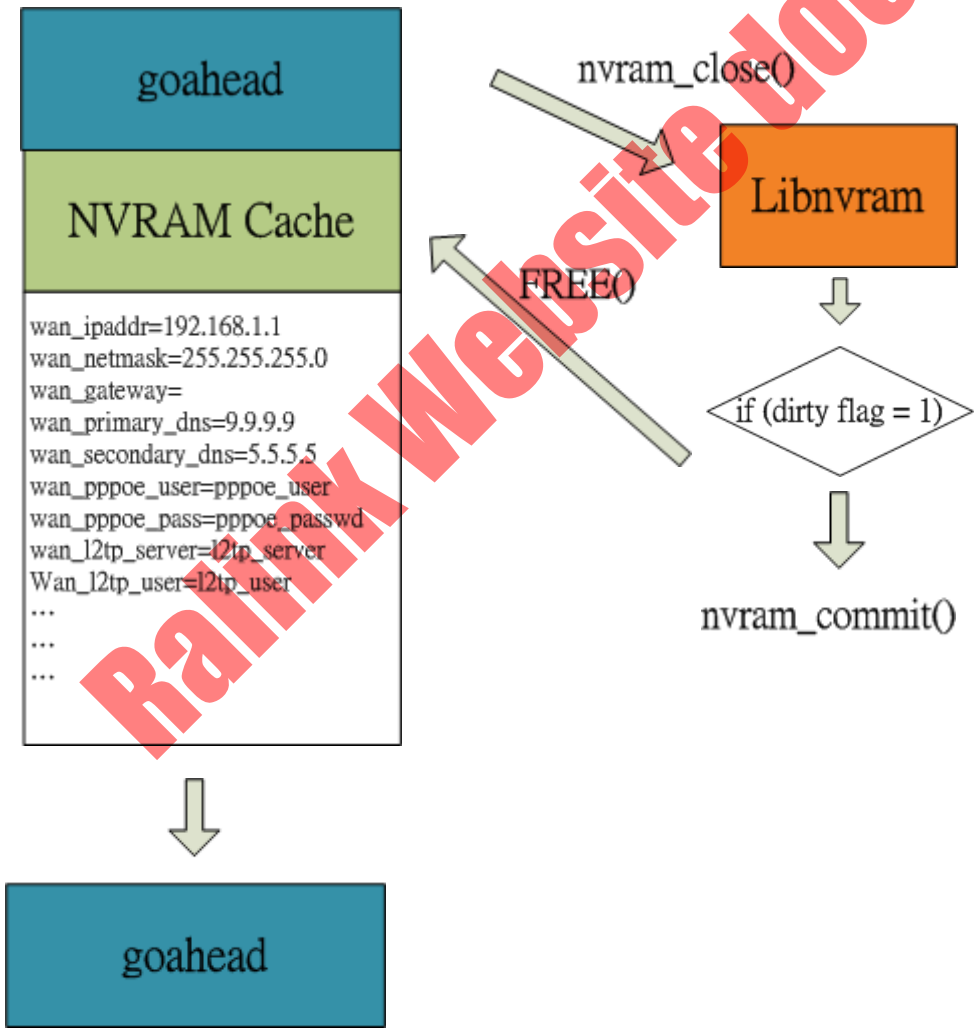

Original Design: nvram_get()

```
char *nvram_get(int index, char *name)
{
        //LIBNV_PRINT("--> nvram_get\n");
        nvram_close(index);
        nvram_init(index);
        return nvram_bufget(index, name);
}
```

- Sync data with Flash and return the freshest value

Original Design: nvram_set()

```
int nvram_set(int index, char *name, char *value)
{
        //LIBNV_PRINT("--> nvram_set\n");
        if (-1 == nvram_bufset(index, name, value))
                return -1;
        return nvram_commit(index);
}
```

- Set the value and commit current data to Flash

Issues?

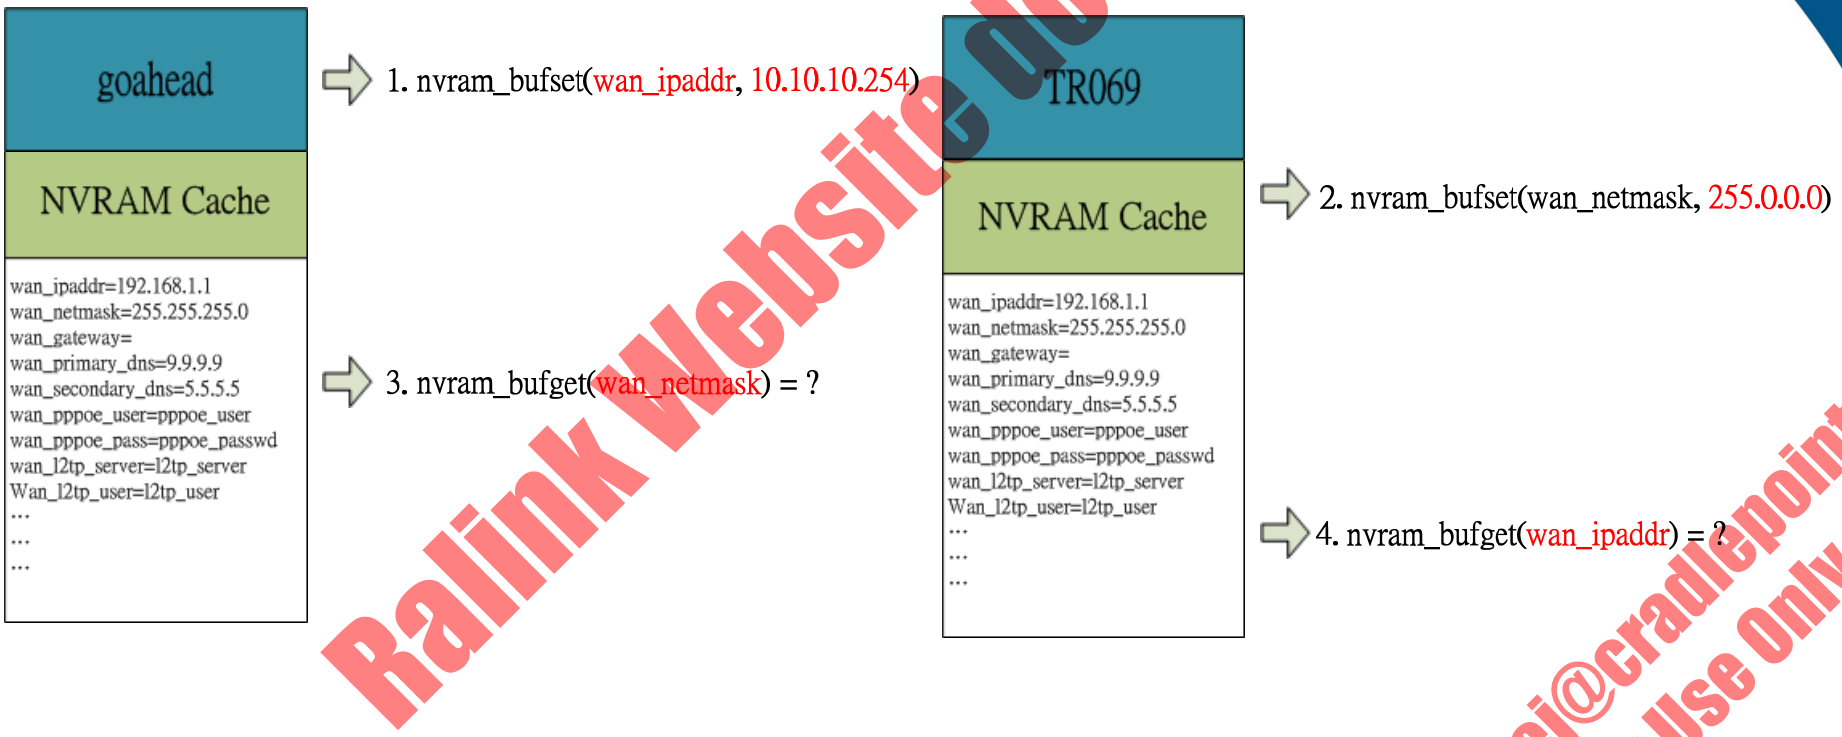

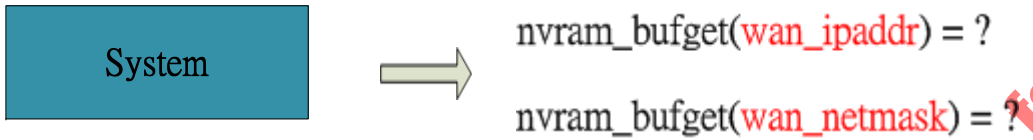

Solutions?

- **Use nvram_get() and nvram_set() instead**
  - nvram_commit() takes long time (flash erase and write)
  - Not efficient
- **Notification**
  - Need to know which program needs to be notified(is using NVRAM)
  - Other program may not be interested in it
- **Centralized Database**
  - Every NVRAM user synchronize changes with it
  - Keep the freshest data

New Design

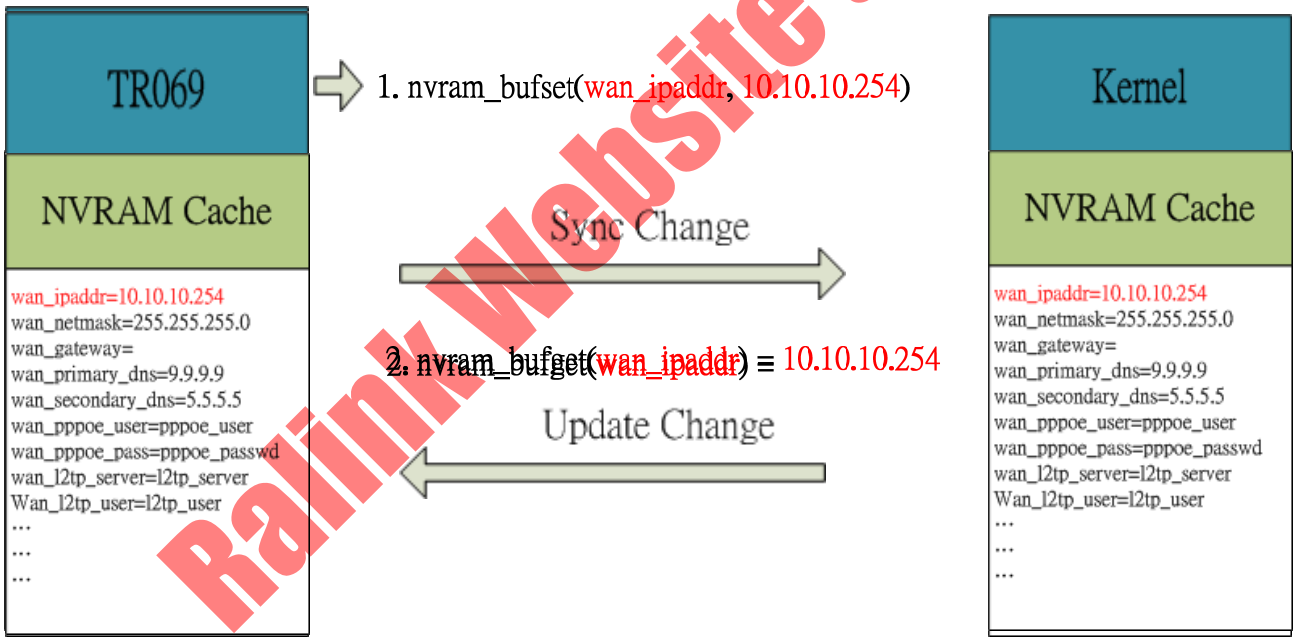

New Design: nvram_init()

- Kernel NVRAM module will invoke nvram_init() at the boot time which does the same thing with previous user-mode nvram_init()

- User-space program also needs to call libnvram nvram_init(), but it gets the NVRAM data from Kernel NVRAM cache instead, not directly from Flash.

New Design:
Kernel nvram_init()

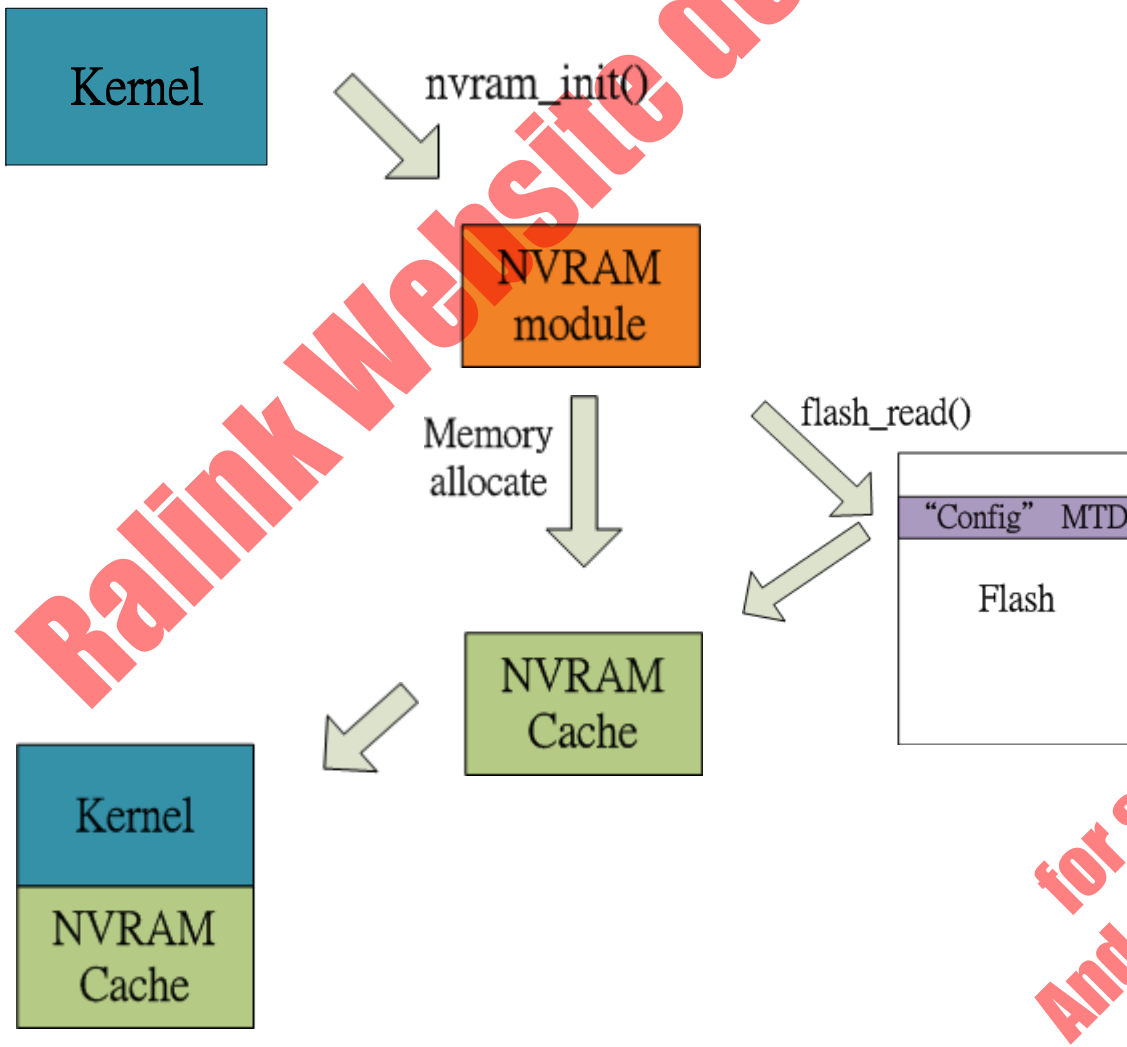

New Design:
Libnvram nvram_init()

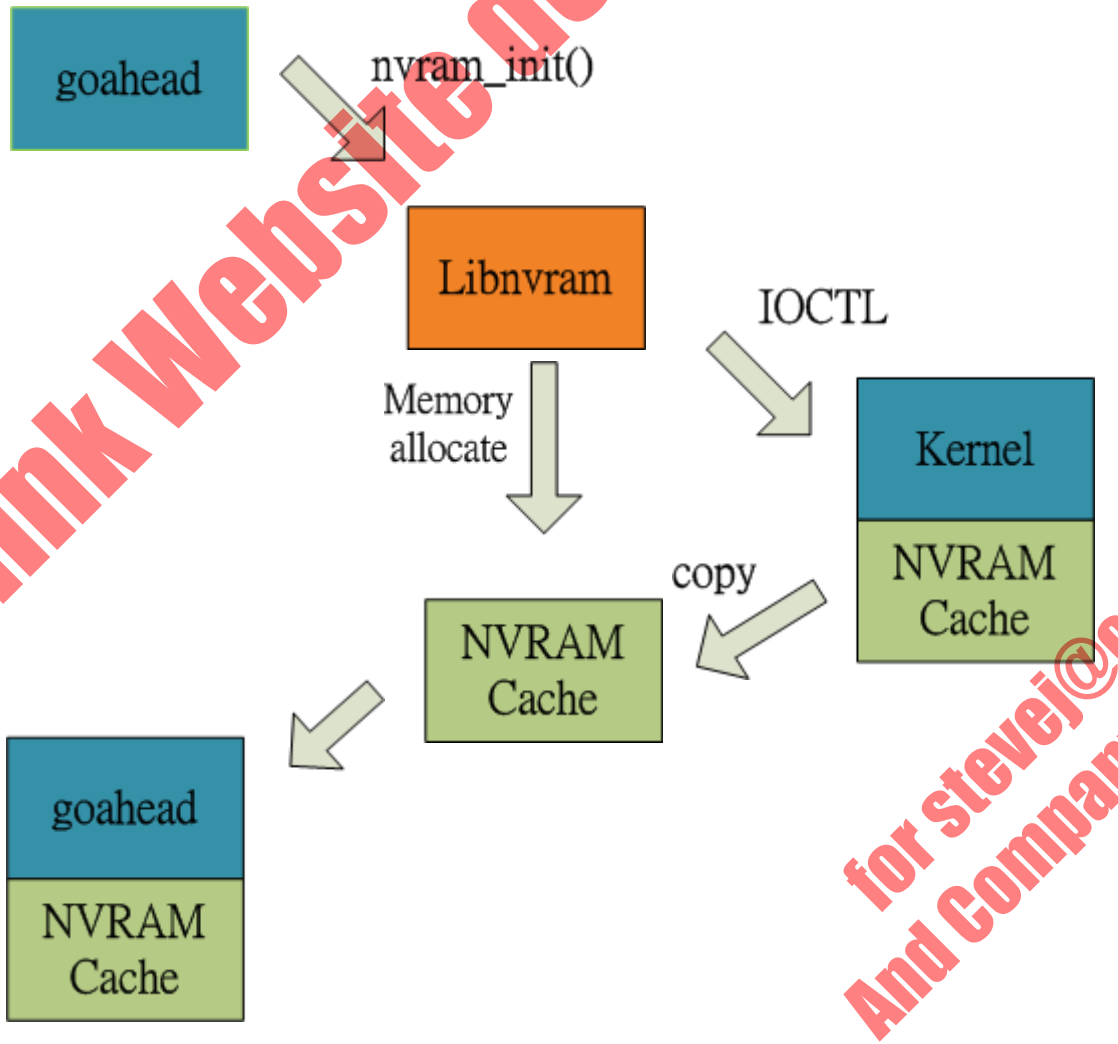

New Design:
Libnvram nvram_get()

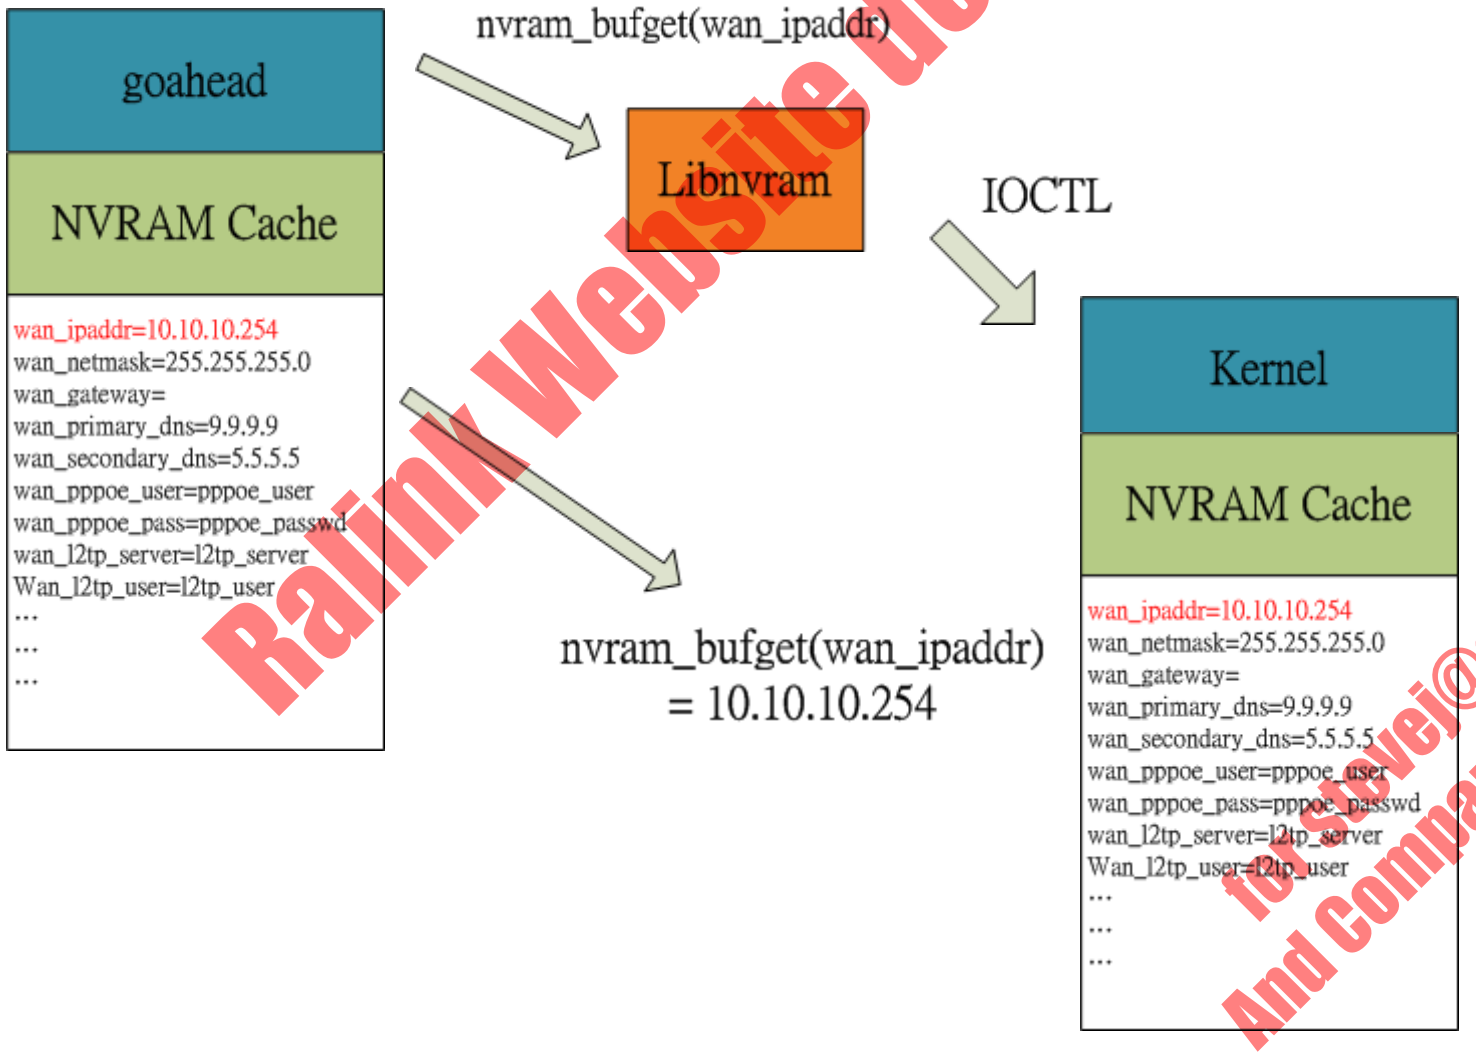

New Design:
Libnvram nvram_set()

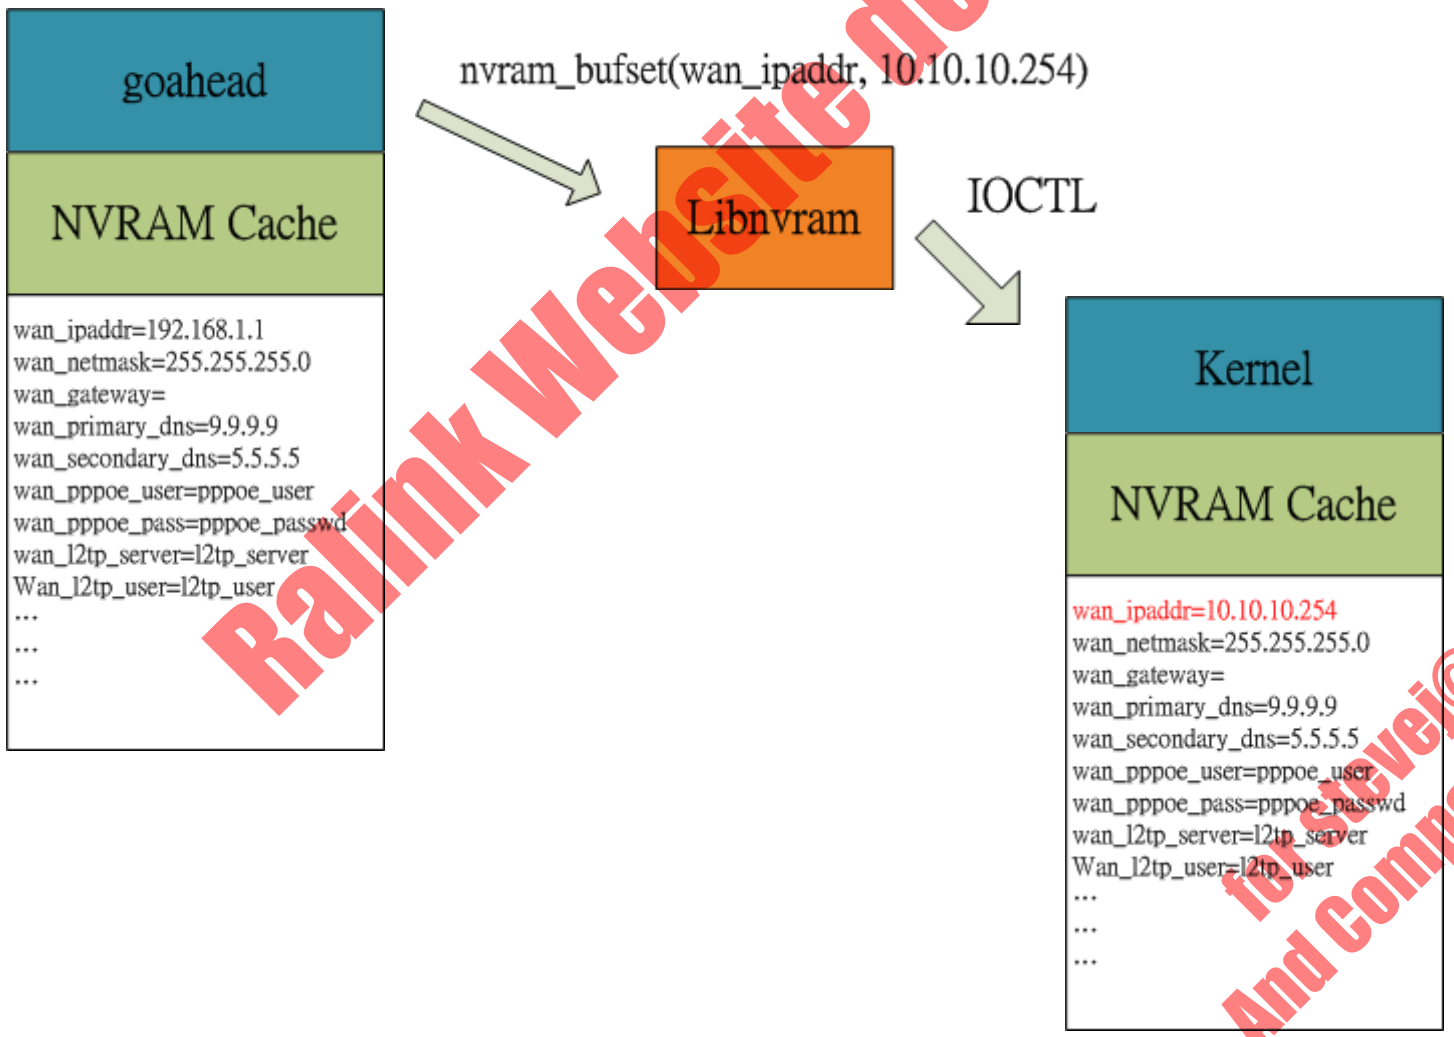

New Design:
Libnvram nvram_commit()

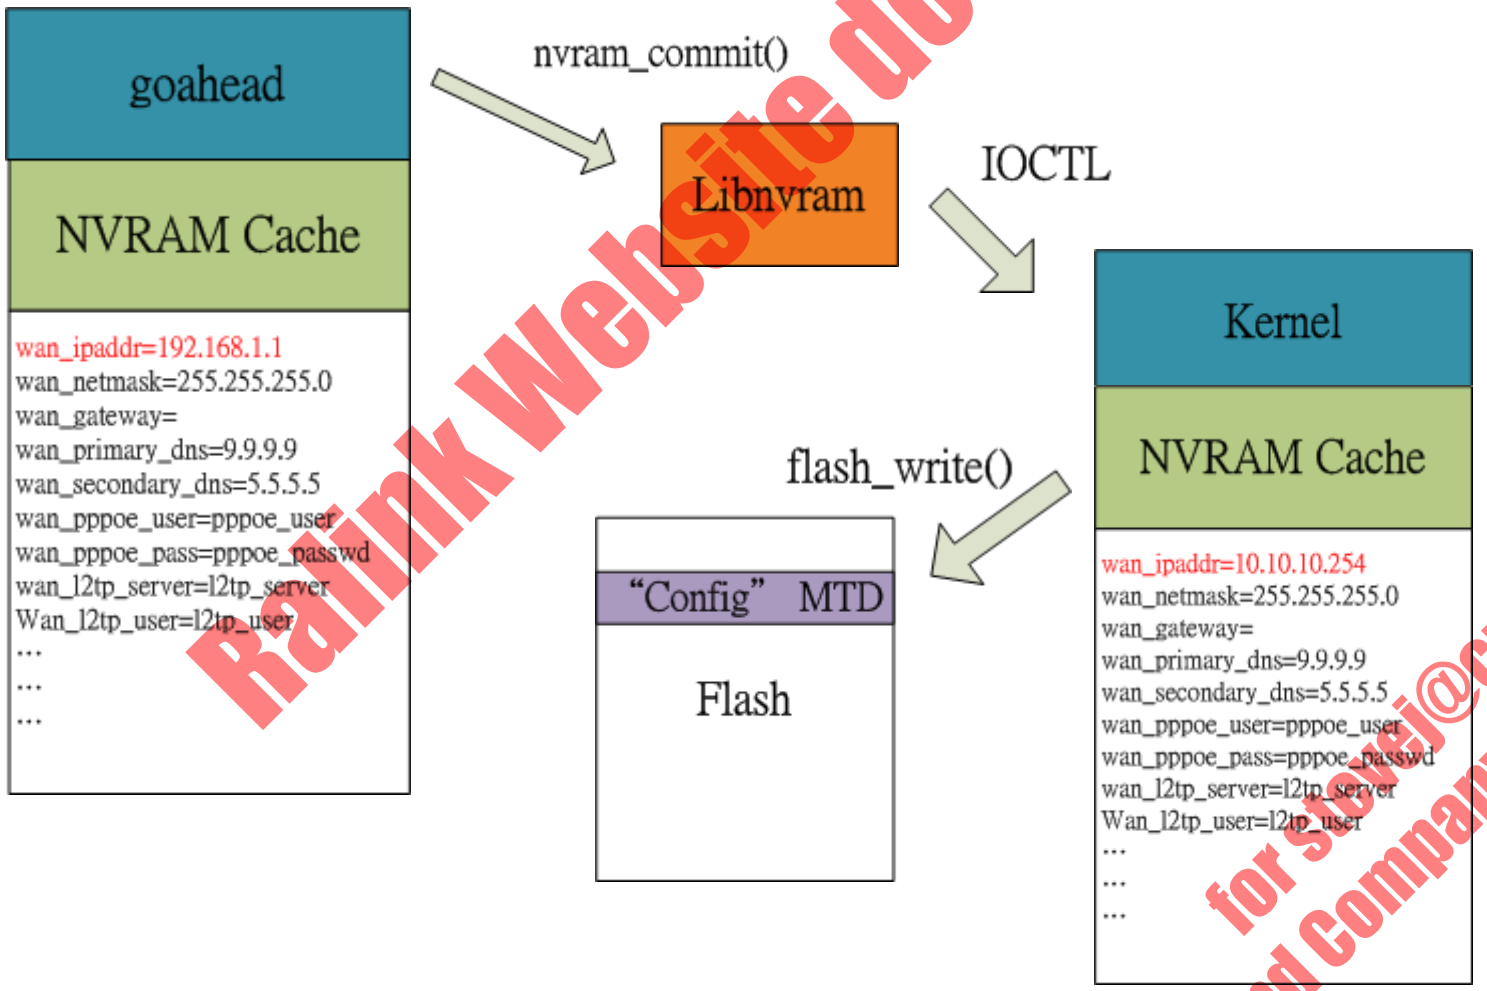

New Design: Libnvram nvram_commit()

- **Which copy of NVRAM cache will be committed to Flash?**
  - ➔ Kernel NVRAM cache
- **Why?**
  - ➔ Kernel NVRAM cache keeps the freshest data

Code Review

- **New added files:**
  - source/linux-2.6.21.x/arch/mips/rt2880/nvram.c
  - source/linux-2.6.21.x/arch/mips/rt2880/nvram.h
  - source/linux-2.6.21.x/arch/mips/rt2880/crc32.c
- **Modified files:**
  - source/lib/libnvram/nvram_env.c
  - source/lib/libnvram/nvram.h
  - source/vendors/Ralink/RT3052/makedevlinks

Test Scenario

- **Memory leak**
  1. Program A keeps calling nvram_bufget()
  2. Program B keeps calling nvram_bufset()
  3. Use “free” command to check if system free memory is decreasing
- **Kernel NVRAM semaphore**
  1. Program A calls X_1 times nvram_bufget(), Y_1 times nvram_bufset(), and Z_1 times nvram_commit()
  2. Program A calls X_2 times nvram_bufget(), Y_2 times nvram_bufset(), and Z_2 times nvram_commit()
  3. Increase global counter in Kernel NVRAM critical setction
  4. Check if the value of global counter = “$X_1 + X_2 + Y_1 + Y_2 + Z_1 + Z_2$ “

Test Scenario

- **Synchronization**
  - Program A and Program B increase a nvram value in turn
  - Program A increase X times, Program B increase Y times
  - Check if that nvram value equals to (X + Y)